

# Unit -IV

Prepared by  
Dr.N.Gobi

# Syllabus- Unit IV

Linux Process Management:

Forms of Process Management-Starting, Pausing, and Resuming  
Processes-Monitoring Processes-Managing Linux Processes-Linux  
Applications: Text EditorsLatex-Email Programs- Network Software-  
Regular Expressions: Grep-Sed-Awk.

# Introduction

- *Process* is a running program.
- Program is static whereas Process is dynamic which changes its states over time.
- In Linux, processes can be started from the GUI or the command line or by other running processes.
- Whenever a process runs, the Linux kernel keeps track of it through a process ID (*PID*).
- Linux kernel is loaded and running, the first process it launches is called system[PID is 1].
- systemd is responsible for starting the run-time environment and then monitoring the environment while in use.

- In Linux, a process can only be created by another process (with the exception of `systemd`). We refer to the creating process as the *parent* and the created process as the *child*.

## Linux System Calls to Create Processes

System Call	Result of the Call
<code>clone()</code>	Like <code>fork</code> (see below) except there is more control over the child process produced with respect to what is duplicated and what is shared between parent and child.
<code>exec()</code>	Take an existing process and replace its image (executable code) with a new image.
<code>fork()</code>	Create a duplicate process of the parent but with its own PID, its own memory and its own resources; parent and child can run concurrently.
<code>vfork()</code>	Same as <code>fork</code> except parent is temporarily suspended and child might be permitted to use the parent's memory space.
<code>wait()</code>	Suspend parent process to wait for an event of a child process.

- fork** is best for independent processes.
- vfork** is faster than fork for processes that immediately exec.
- exec** is for transforming a process into a new program.
- clone** is ideal for creating threads with controlled resource sharing.

# FORMS OF PROCESS MANAGEMENT

- Consider that a program is a piece of software. It contains code and has memory space reserved for data storage.
- Most of this storage is dedicated to variables that are declared and used in the program.
- Program exists in one of two states: the original source code and the executable code. In either case, the program is a *static* entity.

- Process is an *active* entity.
- As the process executes, the values it stores in memory (variables) change.
- As the process executes, it may request resources from the operating system which it then holds, uses and returns to the operating system.
- CPU stores information about the running process via its registers.
- Program Counter (PC) stores the address of the next instruction to fetch from memory while the Instruction register (IR) stores the current instruction.

- Results of the most recent operation are indicated using status flags in the Status Flag (SF) register .
- A run-time stack is maintained for this process, indicating subroutine invocation and parameter passing. The top of the stack is pointed to by the Stack Pointer (SP).



## Process Control Block Information

Item	Use
Process state	The run-time state of the process; usually one of new, ready, running, suspended, terminating, waiting.
Process ID	Usually a numeric designator to differentiate the process from others.
Other process data	Parent process (if there is one), process owner, priority and/or scheduling information, accounting data such as amount of CPU time elapsed.
Process location	Which queue the process currently resides in (ready queue, wait queue, job queue).
Process privilege/state	What mode the process runs in (user, privileged, some other).
Hardware-stored values	PC, IR, SF, SP (and possibly data register values) and interrupt masks.
Resource allocation	I/O devices currently allocated to the process; memory in use; page table.

# *Process Management.*

- Handling the startup and termination of, scheduling and movement of, and interprocess communication of processes.

# Single-process execution

- *single tasking*, the operating system starts a process and then the CPU's attention is entirely held by that one running process.
- CPU executes this one process until either the process terminates or requests some operation by the operating system.
- Among them, the user is limited to running one task at a time and so there is no way to switch back and forth between multiple tasks.
- if the process has to perform some time-consuming input or output, the CPU remains idle while the input/output (I/O) takes place.



# Batch Processing

- *Batch processing* is much like single tasking .Batch processing, the computer is shared among multiple users such that users can submit processes to run at any time.
- Scheduling may be performed by some offline system
- Job may run any time from immediately to hours or days later, the job cannot be expected to run interactively.
- Scheduling algorithm can be used to schedule the processes.

# Concurrent processing

- One of the great inefficiencies of both single tasking and batch processing is that of I/O.
- One process/job may wait for the i/o – time consuming if read from magnetic tapes.
- Idea of switching from a waiting process to a process ready for execution leads us to an improved form of process management called *multiprogramming*.
- In multiprogramming, the CPU executes a single process, but if that process needs to perform time-consuming I/O then that process is set aside and another process is executed.
- Concurrent Processing- By switching between processes in this way, it appears to the user that the computer is executing multiple processes at a time
- Context switching plays a major role.

# For context switching, OS follows

- Pre-emptive
- Non-Preemptive
- Time sharing

# Interrupt handling

- Process of interrupting the CPU as an *interrupt* and the process of requesting an interrupt as an *interrupt request* (IRQ).
- IRQ may originate from hardware or software. For hardware, the IRQ is carried over a reserved line on the bus connecting the hardware device to the CPU (or to an interrupt controller device). For software, an IRQ is submitted as an interrupt signal.
- Upon receiving an interrupt request, the CPU finishes the current fetch-execute cycle and then decides how to respond to the interrupt request.
- For every type of interrupt, the operating system contains an interrupt handler.
- Each *interrupt handler* is a piece of code written to handle a specific type of interrupting situation.



# STARTING, PAUSING AND RESUMING PROCESSES

- Depending on the version of Linux and the GUI being used, there are menu selections and icons.
- Parent process making the request might be the GUI, Bash, another running application or a part of the operating system (e.g., a service).
- kernel's process manager creates the process and prepares it for execution.
- When a process is launched by the user, the process runs under access rights based on its *Effective User ID* (EUID) and *Effective Group ID* (EGID).

# Launching processes from the command line

- Launching processes from within a shell is accomplished by specifying the program name. → Absolute path or relative path.

# Suspending and resuming processes from the command line

- Run as many processes in the background.
- A terminal window can only have one process running in the foreground at a time.
- Foreground process is the one that provides interaction with the user, whether to receive user input or to output information to the window, or both.
- Jobs command responds with all of the active jobs in the terminal window. These are processes that have been started from the command line but have not yet terminated.

- A process may be in a *stopped* state. This means that the process has been running but is currently suspended.
- User can resume a suspended process. In doing so, we specify whether to resume the process in the foreground or in the background.
- In order to stop a running, foreground process, **type control+z in the terminal window**

# MONITORING PROCESSES

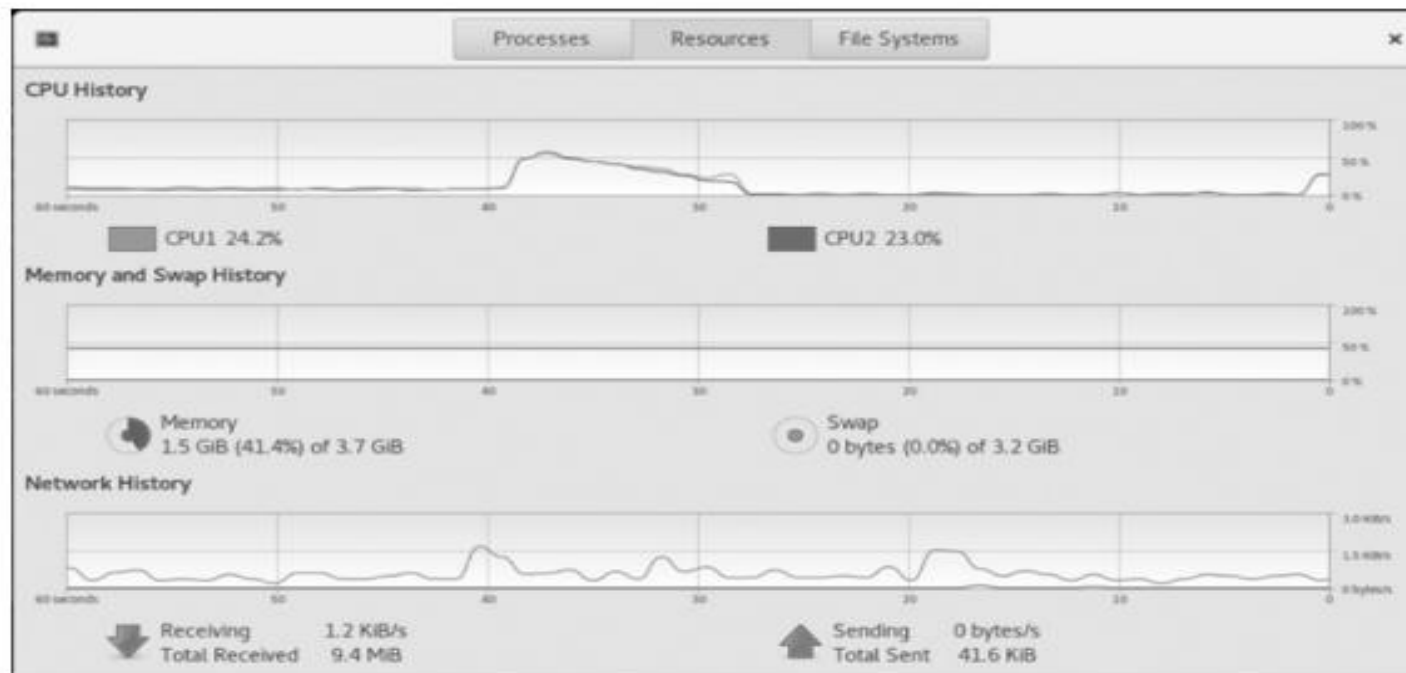
- Operating system can run multiple processes efficiently with little or no user intervention.
- Several different tools available to monitor processes and system resources.
- Primary GUI tool is called the System Monitor.
- Launching System Monitor from the GUI has changed.
- Two different versions of the System Monitor, one for Gnome and one for KDE (an alternative GUI desktop).

# Gnome

Processes									
Process Name	User	% CPU	ID	Memory	Disk read tota	Disk write tot	Disk read	Disk write	Priority
gnome-shell	foxr	4	2285	197.4 MB	8.2 MiB	852.0 KiB	N/A	N/A	Normal
gnome-system-monitor	foxr	1	6279	14.4 MB	792.0 KiB	N/A	N/A	N/A	Normal
systemd	foxr	0	2193	1.7 MB	161.0 KiB	16.0 KiB	N/A	N/A	Normal
(sd-pam)	foxr	0	2199	5.2 MB	N/A	N/A	N/A	N/A	Normal
pulseaudio	foxr	0	2212	1.5 MB	616.0 KiB	8.0 KiB	N/A	N/A	Very High
gnome-keyring-daemon	foxr	0	2218	1.1 MB	N/A	N/A	N/A	N/A	Normal
dbus-daemon	foxr	0	2225	1.5 MB	N/A	N/A	N/A	N/A	Normal
gdm-wayland-session	foxr	0	2228	1.6 MB	4.0 KiB	N/A	N/A	N/A	Normal
gnome-session-binary	foxr	0	2231	2.9 MB	36.0 MiB	8.0 KiB	N/A	N/A	Normal
gvfsd	foxr	0	2308	1.1 MB	316.0 KiB	N/A	N/A	N/A	Normal
gvfsd-fuse	foxr	0	2317	3.1 MB	356.0 KiB	N/A	N/A	N/A	Normal
Xwayland	foxr	0	2327	10.5 MB	N/A	24.0 KiB	N/A	N/A	Normal
at-spi-bus-launcher	foxr	0	2333	1.0 MB	N/A	N/A	N/A	N/A	Normal
dbus-daemon	foxr	0	2338	544.0 KiB	N/A	N/A	N/A	N/A	Normal
at-spi2-registr	foxr	0	2343	1012.0 KiB	N/A	N/A	N/A	N/A	Normal
ibus-dconf	foxr	0	2350	812.0 KiB	24.0 KiB	N/A	N/A	N/A	Normal
ibus-extension-att	foxr	0	2353	7.5 MB	600.0 KiB	N/A	N/A	N/A	Normal

End Process

- Resources tab provides a summary of CPU, memory/swap space and network utilization over time (the last minute).



- File Systems tab displays statistics about the file system.



# Command-Line Monitoring Tools

- System Monitor provides a convenient way to view system usage and control running processes, we might be more interested in exploring system usage from the command line for two reasons.
- similar to the idea of opening multiple windows to run multiple processes that we explored in the last section, using the System Monitor takes more resources than command-line programs.
- information we can get from the System Monitor is not as comprehensive as what can be found from some of the command-line programs. So, we turn to two of the primarily monitoring tools:  
**top and ps**

# top program

- **top** program is launched from the command line but unlike most of the commands we have viewed, this command fills the terminal window with its output and remains running.
- **top** is an interactive program that updates its output in a specified refresh rate and changes its appearance based on input keystrokes.
- **Default refresh rate is 3 seconds.**

File Edit View Search Terminal Help

```
top - 09:22:02 up 1 day, 20:34, 1 user, load average: 0.06, 0.02, 0.00
Tasks: 290 total, 3 running, 283 sleeping, 4 stopped, 0 zombie
%Cpu(s): 0.2 us, 0.0 sy, 0.0 ni, 99.6 id, 0.0 wa, 0.0 hi, 0.3 si, 0.0 st
MiB Mem : 1789.5 total, 79.3 free, 1107.6 used, 602.6 buff/cache
MiB Swap: 2048.0 total, 1623.4 free, 424.6 used. 504.4 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
944	root	20	0	575544	6772	5948	S	0.3	0.4	67:03.88	vmtoolsd
2299	foxr	20	0	4361240	162476	40476	S	0.3	8.9	1:33.92	gnome-s+
27078	root	20	0	0	0	0	I	0.3	0.0	0:03.92	kworker+
27122	foxr	20	0	275272	5228	4368	R	0.3	0.3	0:00.15	top
1	root	20	0	245520	6640	4464	S	0.0	0.4	0:05.05	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.10	kthreadd
3	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_gp
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_par+
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.01	kworker+
9	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_perc+
10	root	20	0	0	0	0	S	0.0	0.0	0:01.07	ksoftir+
11	root	20	0	0	0	0	I	0.0	0.0	0:06.45	rcu_sch+
12	root	rt	0	0	0	0	S	0.0	0.0	0:00.07	migrati+
13	root	rt	0	0	0	0	S	0.0	0.0	0:00.01	watchdo+
14	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/0
15	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/1
16	root	rt	0	0	0	0	S	0.0	0.0	0:00.03	watchdo+

# ps command

- ps command provides a detailed examination of the running processes as a snapshot.
- Unlike top, it is not interactive. It displays the information and exits.

# MANAGING PROCESS PRIORITY

- Linux runs processes and threads together by switching off between them.
- Because of the speed and power of modern processors, the user is typically unaware of the time elapsing as the processor moves from one process to another and back.
- processor moves through the processes in the ready queue and back to the first in under a millisecond (thousandths of a second).

- A process is a standalone entity; it has its own code, data and status information.
- Thread, some-times called a *lightweight* process, shares its code and data (or at least some data) with other threads.
- Threads are, in essence, portions of a process.
- Linux executes processes and threads using multitasking and multithreading.
- Linux command batch, for instance, forces a process to run in batch mode.
- A batch process is one that does not interact with the user. Therefore, any batch process must be provided its input at the time the process is launched.

- In Linux, a process' priority is established by setting its *niceness* value. Niceness refers to how nice a process is with respect to other processes.
- With a higher niceness value, a process will offer some of its CPU time to other processes.
- Higher the niceness value, the less CPU time it will take when it is that process' turn, and thus higher niceness means lower priority.

<https://www.geeksforgeeks.org/priority-of-process-in-linux-nice-value/>

gnome-shell	fovr	17	2285	197.4 MiB	8.2 MiB	968.0 KiB	N/A
gnome-system-monitor	fovr				4.0 KiB	N/A	1.3 KiB/s
systemd	fovr				161.0 KiB	16.0 KiB	N/A
(sd-pam)	fovr				N/A	N/A	N/A
pulseaudio	fovr					8.0 KiB	N/A
gnome-keyring-daemon	fovr					N/A	N/A
dbus-daemon	fovr					N/A	N/A
gdm-wayland-session	fovr					N/A	N/A
gnome-session-binary	fovr					8.0 KiB	N/A
gvfsd	fovr	0	2308	1.1 MiB		N/A	N/A
gvfsd-fuse	fovr	0	2317	3.1 MiB		N/A	N/A

Properties

Alt+Return

Memory Maps

Ctrl+M

Open Files

Ctrl+O

Change Priority

☐ Very High

Stop

Ctrl+S

☐ High

Continue

Ctrl+C

☒ Normal

End

Ctrl+E

☐ Low

Kill

Ctrl+K

☐ Very Low☐ Custom



# Regular Expressions

- **Linux Regular Expressions** are special characters which help search data and matching complex patterns.
- Regular expressions are shortened as 'regexp' or 'regex'.
- They are used in many Linux programs like grep, bash, rename, sed, etc.

# Basic Regular expressions

Symbol	Descriptions
.	replaces any character
^	matches start of string
\$	matches end of string
*	matches up zero or more times the preceding character
\	Represent special characters
()	Groups regular expressions
?	Matches up exactly one character

# Syntax:

- `cat sample | grep string`

```
guru99@guru99-VirtualBox:~$ cat sample
```

```
apple
```

```
bat
```

```
ball
```

```
ant
```

```
eat
```

```
pant
```

```
people
```

```
taste
```

```
guru99@guru99-VirtualBox:~$
```

```
guru99@guru99-VirtualBox:~$ cat sample | grep a
apple
bat
ball
ant
eat
pant
taste
guru99@guru99-VirtualBox:~$
```

```
guru99@guru99-VirtualBox:~$ cat sample | grep ^a
apple
ant
guru99@guru99-VirtualBox:~$
```

```
guru99@guru99-VirtualBox:~$ cat sample | grep t
```

bat

ant

eat

pant

taste

```
guru99@guru99-VirtualBox:~$ cat sample | grep t$
```

bat

ant

eat

pant

```
guru99@guru99-VirtualBox:~$
```

# Interval Regular expressions

- These expressions tell us about the number of occurrences of a character in a string.

Expression	Description
$\{n\}$	Matches the preceding character appearing 'n' times exactly
$\{n,m\}$	Matches the preceding character appearing 'n' times but not more than m
$\{n, \}$	Matches the preceding character only when it appears 'n' times or more

Filter out all lines that contain character 'p'

```
guru99@guru99-VirtualBox:~$ cat sample|grep p  
apple  
pant  
people
```



Check that the character 'p' appears exactly 2 times in a string one after the other.

```
guru99@guru99-VirtualBox:~$ cat sample|grep -E p\{2}
apple
guru99@guru99-VirtualBox:~$
```

need to add -E with these regular expressions.

# Extended regular expressions

- regular expressions contain combinations of more than one expression.

Expression	Description
\+	Matches one or more occurrence of the previous character
\?	Matches zero or one occurrence of the previous character

```
guru99@guru99-VirtualBox:~$ cat sample|grep t
bat
ant
eat
pant
taste
```

Filter out lines where character 'a' precedes character 't'

```
guru99@guru99-VirtualBox:~$ cat sample|grep "a\+t"
bat
eat
guru99@guru99-VirtualBox:~$
```

# Brace expansion

- Brace expansion is either a sequence or a comma separated list of items inside curly braces “{}”.
- Starting and ending items in a sequence are separated by two periods “..”.

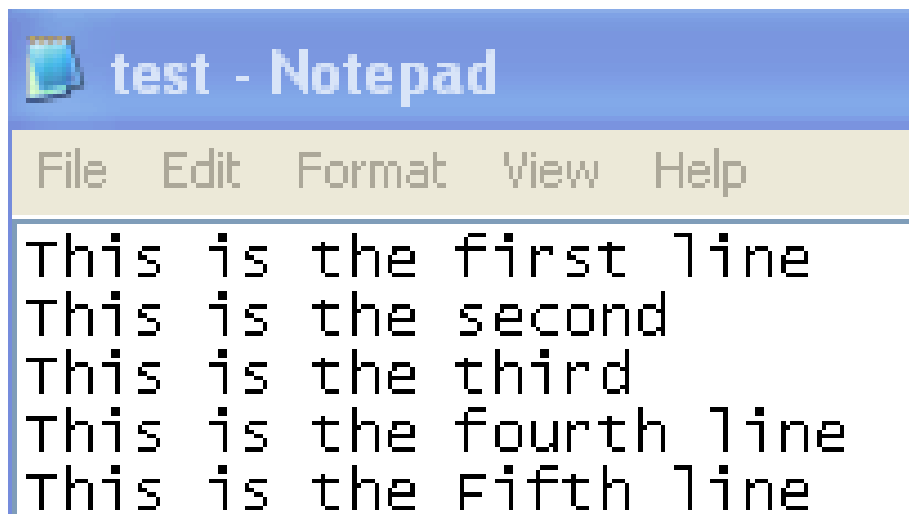
```
guru99@VirtualBox:~$ echo {aa,bb,cc,dd}
aa bb cc dd
guru99@VirtualBox:~$ echo {0..11}
0 1 2 3 4 5 6 7 8 9 10 11
guru99@VirtualBox:~$ echo {a..z}
a b c d e f g h i j k l m n o p q r s t u v w x y z
guru99@VirtualBox:~$ echo a{0..9}b
a0b a1b a2b a3b a4b a5b a6b a7b a8b a9b
guru99@VirtualBox:~$
```

# Matching Control

- `'-i'`
  - Ignores case.
- `'-v'`
  - Inverts the matching. When used, grep will print out lines that do not match the pattern
- `'-e pattern'`
  - Pattern is the pattern. This can be used to specify multiple patterns, or if the pattern starts with a `'-'`. A line only has to contain one of the patterns to be matched.

# Examples using '-i', '-v'

- The file below was used for these



```
test - Notepad
File Edit Format View Help
This is the first line
This is the second
This is the third
This is the fourth line
This is the Fifth line
```

```
Colin@masterson1 /cygdrive/c/Documents a
h_year/cs265
$ grep -i 'LINE' test.txt
This is the first line
This is the fourth line
This is the Fifth line

Colin@masterson1 /cygdrive/c/Documents a
h_year/cs265
$ grep -v 'line' test.txt
This is the second
This is the third

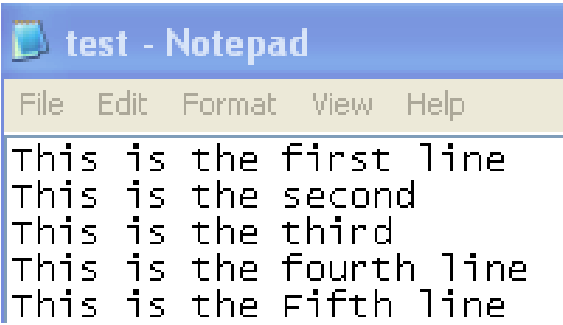
Colin@masterson1 /cygdrive/c/Documents a
h_year/cs265
$ grep -e 'fourth' -e 'first' test.txt
This is the first line
This is the fourth line
```

# General Output Control

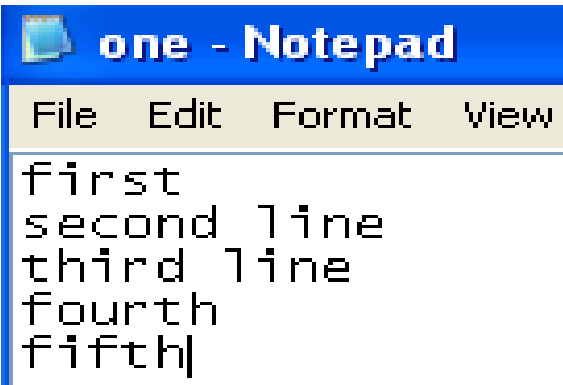
- `'-c'`
  - Suppress normal output and instead print out a count of matching lines for each input file
- `'-l'`
  - Suppress normal output and print the name of each file that contains at least one match
- `'-L'`
  - Suppress normal output. Print the name of each file that does not contain any matches
- Note: both the `'-l'` and `'-L'` options will stop searching a file once a match is found

# Examples using '-c', '-l', and '-L'

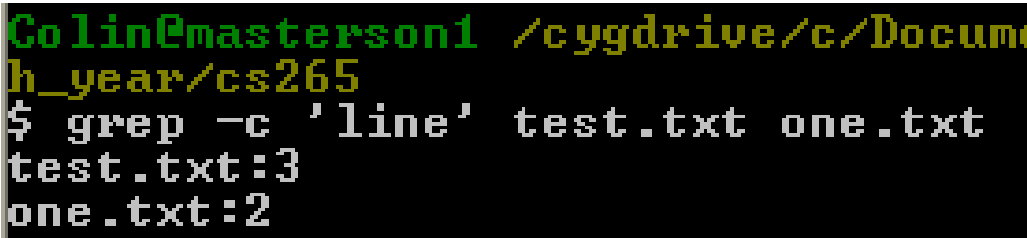
- Two files were used and specified below




```
test - Notepad
File Edit Format View Help
This is the first line
This is the second
This is the third
This is the fourth line
This is the Fifth line
```



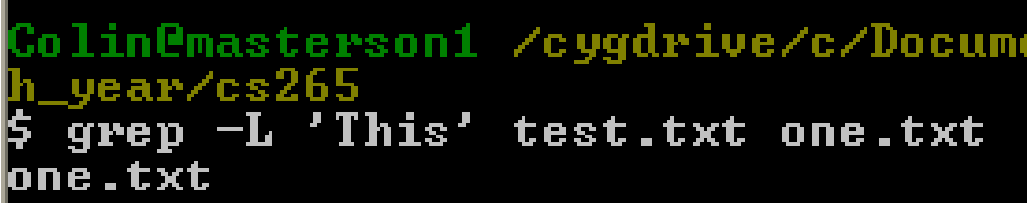
```
one - Notepad
File Edit Format View
first
second line
third line
fourth
fifth|
```



```
Colin@masterson1 /cygdrive/c/Documents and Settings/h_year/cs265
$ grep -c 'line' test.txt one.txt
test.txt:3
one.txt:2
```



```
Colin@masterson1 /cygdrive/c/Documents and Settings/h_year/cs265
$ grep -l 'This' test.txt one.txt
test.txt
```



```
Colin@masterson1 /cygdrive/c/Documents and Settings/h_year/cs265
$ grep -L 'This' test.txt one.txt
one.txt
```

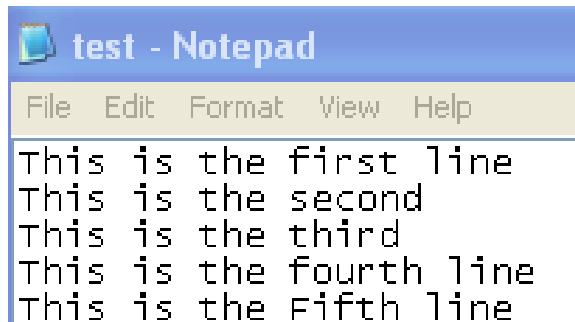


# Output Line Prefix Control

- `'-n'`
  - Prefixes each line of output with the line number from the input file the match was found on
- `'-H'`
  - Prefix each line of output with the input file name that the match was found in
- `'-T'`
  - Makes sure that the actual line content (or whatever content comes after the `'-T'`) lands on a tab stop

# Examples using '-H', '-n', and '-T'

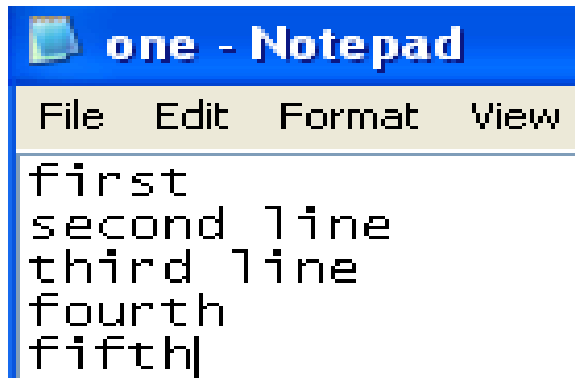
- Two files were used and specified below



test - Notepad

File Edit Format View Help

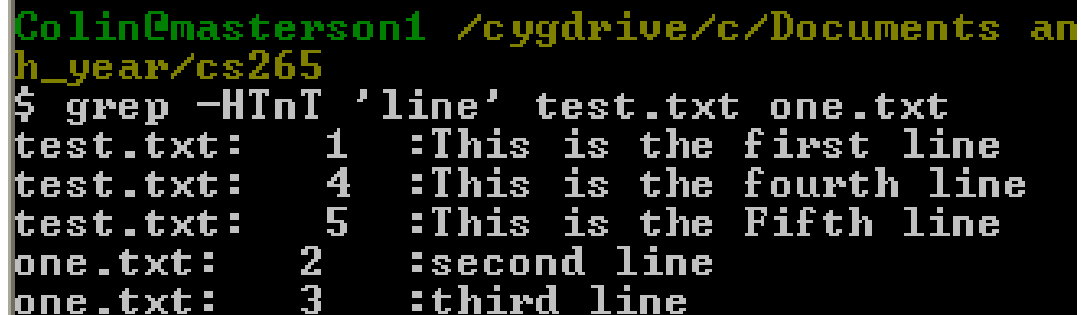
This is the first line  
This is the second  
This is the third  
This is the fourth line  
This is the Fifth line



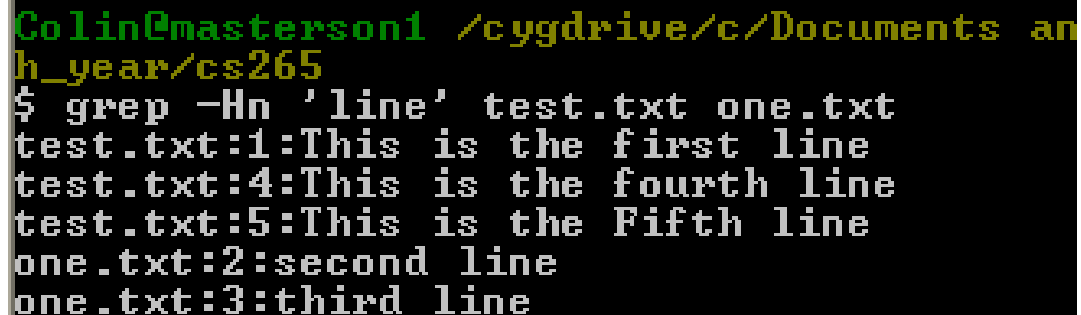
one - Notepad

File Edit Format View

first  
second line  
third line  
fourth  
fifth|



```
Colin@masterson1 /cygdrive/c/Documents and Settings/h_year/cs265
$ grep -HTnT 'line' test.txt one.txt
test.txt: 1 :This is the first line
test.txt: 4 :This is the fourth line
test.txt: 5 :This is the Fifth line
one.txt: 2 :second line
one.txt: 3 :third line
```



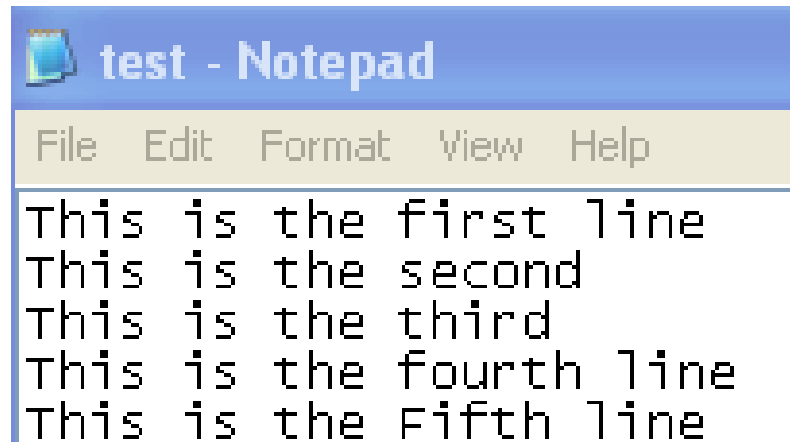
```
Colin@masterson1 /cygdrive/c/Documents and Settings/h_year/cs265
$ grep -Hn 'line' test.txt one.txt
test.txt:1:This is the first line
test.txt:4:This is the fourth line
test.txt:5:This is the Fifth line
one.txt:2:second line
one.txt:3:third line
```

# Context Line Control

- `'-A num'`
  - Print *num* lines of trailing context after matching lines
- `'-B num'`
  - Print *num* lines of leading context before matching lines
- `'-C num'` or `'-num'`
  - Print *num* lines of leading and trailing output context

# Examples using '-A',

- The file below was used for the



```
test - Notepad
File Edit Format View Help
This is the first line
This is the second
This is the third
This is the fourth line
This is the Fifth line
```

```
Colin@masteron1 /cygdrive/c
h_year/cs265
$ grep -A 1 'third' test.txt
This is the third
This is the fourth line
```

```
Colin@masteron1 /cygdrive/c
h_year/cs265
$ grep -B 1 'third' test.txt
This is the second
This is the third
```

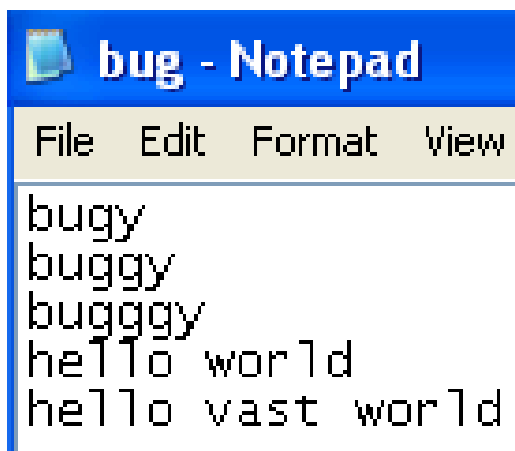
```
Colin@masteron1 /cygdrive/c
h_year/cs265
$ grep -C 1 'third' test.txt
This is the second
This is the third
This is the fourth line
```

```
Colin@masteron1 /cygdrive/c
h_year/cs265
$ grep -l 'third' test.txt
This is the second
This is the third
This is the fourth line
```

# Special Characters

- '.' The period '.' matches any single character.
- '?' The preceding item is optional and will be matched at most once.
- '\*' The preceding item will be matched zero or more times.
- '+' The preceding item will be matched one or more times.
- '{n}' The preceding item is matched exactly n times.
- '{n,}' The preceding item is matched n or more times.
- '{,m}' The preceding item is matched at most m times.
- '{n,m}' The preceding item is matched at least n times, but not more than m times.

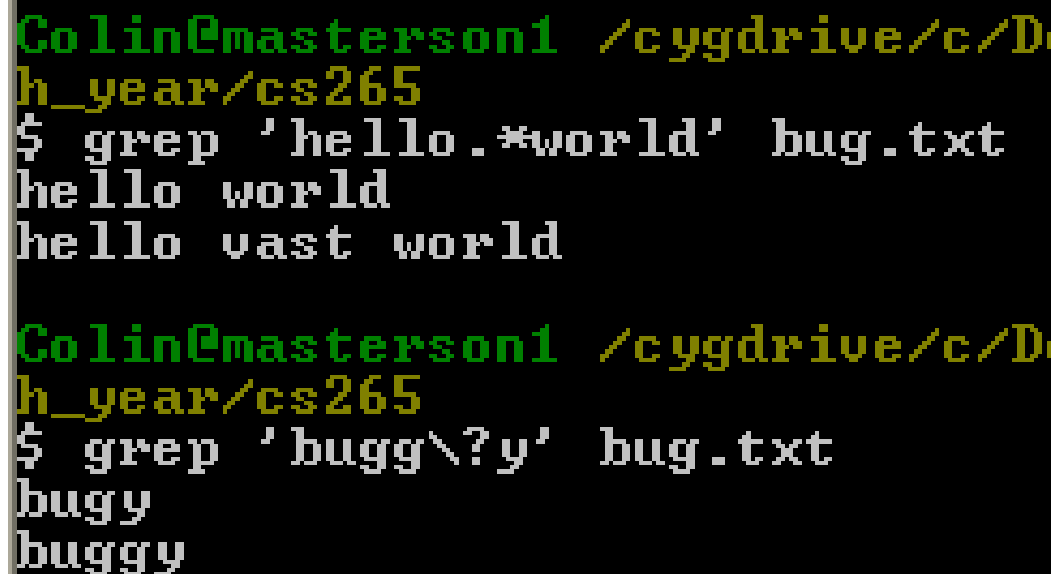
# Examples using '.', '\*', and '?'



bug - Notepad

File Edit Format View

bugy  
buggy  
bugggy  
hello world  
hello vast world



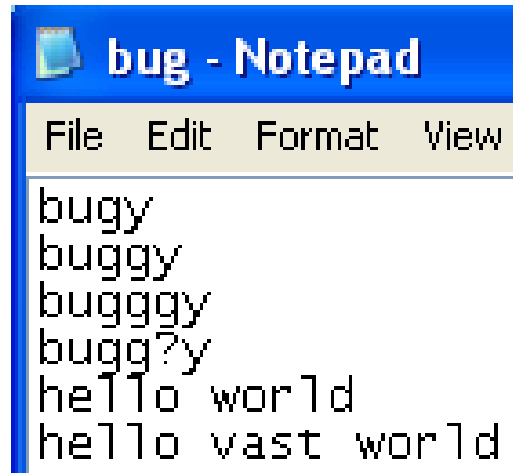
```
Colin@masterson1 /cygdrive/c/D
h_year/cs265
$ grep 'hello.*world' bug.txt
hello world
hello vast world

Colin@masterson1 /cygdrive/c/D
h_year/cs265
$ grep 'bugg\?y' bug.txt
bugy
buggy
```

# Basic vs Extended Regular Expressions

- `'-G'`
  - Interpret pattern as basic regular expression (BRE). This is the default.
- `'-E'`
  - Interpret pattern as extended regular expression (ERE)
- When using basic regular expression some special characters (like `'?'` in the previous example) lose their special meaning and must have a `'\'`, the escape character, before them
- When using ERE, the escape character is unnecessary

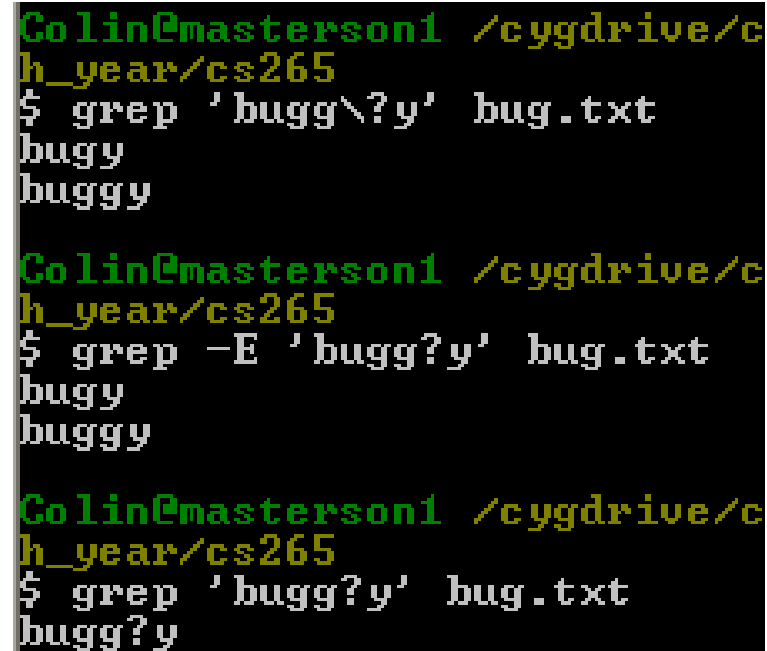
# BRE and ERE Difference



bug - Notepad

File Edit Format View

bugy  
buggy  
bugggy  
bugg?y  
hello world  
hello vast world



```
Colin@masterson1 /cygdrive/c
h_year/cs265
$ grep 'bugg\?y' bug.txt
bugy
buggy

Colin@masterson1 /cygdrive/c
h_year/cs265
$ grep -E 'bugg?y' bug.txt
bugy
buggy

Colin@masterson1 /cygdrive/c
h_year/cs265
$ grep 'bugg?y' bug.txt
bugg?y
```

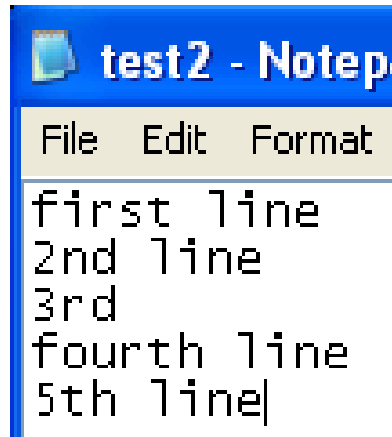
- Note that without the `\` in the BRE call (example 3), the `'?'` is seen as a normal character



# Bracket Expressions

- A bracket expression is a list of characters enclosed by '[' and ']'. It matches any single character in the list
- However, if the first character in the list is '^', it matches any character not in the list
- A range can be done by using '-' in a bracket expression
  - [0-5] is the same as [012345]
- Some ranges are pre-defined in character classes
  - [:digit:] is the same as 0123456789
  - When using grep, the class name (including brackets) must be contained within another set of brackets

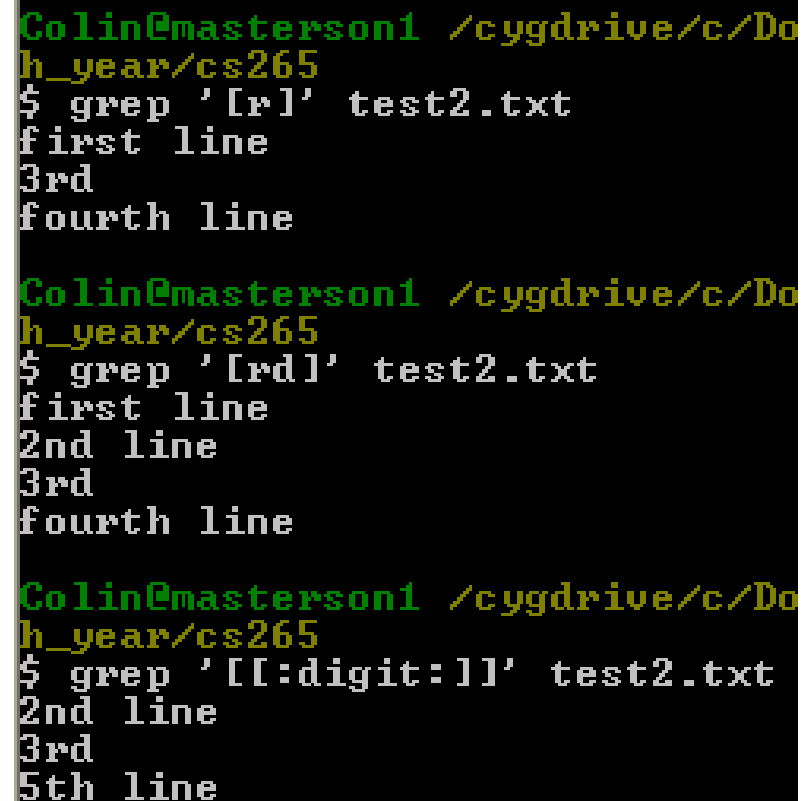
# Bracket Expression Example



test2 - Notepad

File Edit Format

first line  
2nd line  
3rd  
fourth line  
5th line



```
Colin@masterson1 /cygdrive/c/Do  
h_year/cs265  
$ grep '[r]' test2.txt  
first line  
3rd  
fourth line  
  
Colin@masterson1 /cygdrive/c/Do  
h_year/cs265  
$ grep '[rd]' test2.txt  
first line  
2nd line  
3rd  
fourth line  
  
Colin@masterson1 /cygdrive/c/Do  
h_year/cs265  
$ grep '[[digit:]]' test2.txt  
2nd line  
3rd  
5th line
```

# awk

- awk is a programmable, pattern-matching, and processing tool available in UNIX.
- It works equally well with text and numbers.
- It derives its name from the first letter of the last name of its three authors namely Alfred V. Aho, Peter J. Weinberger and Brian W. Kernighan.

# Simple awk Filtering

- awk is not just a command, but a programming language too.
- awk utility is a pattern scanning and processing language.
- It searches one or more files to see if they contain lines that match specified patterns and then perform associated actions, such as writing the line to the standard output or incrementing a counter each time it finds a match.

# AWK Operations:

- (a) Scans a file line by line
- (b) Splits each input line into fields
- (c) Compares input line/fields to pattern
- (d) Performs action(s) on matched lines

```
awk 'pattern { action }' filename
```

### Common Patterns and Special Variables:

- **BEGIN** - Executes before processing any lines
- **END** - Executes after processing all lines
- **NR** - Current line number
- **NF** - Number of fields in current line
- **\$0** - Entire line
- **\$1, \$2, ...** - First field, second field, etc.
- **FS** - Field separator (default is whitespace)
- **OFS** - Output field separator

- **NR:** NR command keeps a current count of the **number of input records**.
- **NF:** NF command keeps a count of the number of fields within the **current input record**.
- **FS:** FS command contains the **field separator** character which is used to divide fields on the input line → default is “white space”, meaning space and tab characters.
- **RS:** RS command stores the current **record separator** character.
- **OFS:** OFS command stores the **output field separator**, which separates the fields when Awk prints them.
- **ORS:** ORS command stores the **output record separator**, which separates the output lines when Awk prints them.

\$cat > employee.txt

ajay manager account 45000

sunil clerk account 25000

varun manager sales 50000

amit manager account 47000

tarun peon sales 15000

deepak clerk sales 23000

sunil peon sales 13000

satvik director purchase 80000



Print the lines which match the given pattern.

```
awk '{print}' employee.txt
```

ajay manager account 45000

sunil clerk account 25000

varun manager sales 50000

amit manager account 47000

tarun peon sales 15000

deepak clerk sales 23000

sunil peon sales 13000

satvik director purchase 80000

```
awk '/manager/ {print}' employee.txt
```

ajay manager account 45000

varun manager sales 50000

amit manager account 47000

# Print entire file

```
awk '{print}' file.txt
```

# or

```
awk '{print $0}' file.txt
```

# Print first column

```
awk '{print $1}' file.txt
```

# Print first and second columns

```
awk '{print $1, $2}' file.txt
```

# Print last column

```
awk '{print $NF}' file.txt
```

# Print second-to-last column

```
awk '{print $(NF-1)}' file.txt
```

# Print line numbers

```
awk '{print NR, $0}' file.txt
```

# Print line numbers for non-empty lines

```
awk 'NF {print NR, $0}' file.txt
```

# Print line numbers and first column

```
awk '{print NR ":", $1}' file.txt
```

# Print lines containing "pattern"

awk '/pattern/' file.txt

# Print lines NOT containing "pattern"

awk '!/pattern/' file.txt

# Print lines longer than 80 characters

awk 'length(\$0) > 80' file.txt

# Print lines with more than 5 fields

awk 'NF > 5' file.txt

# Print number of fields in each line

```
awk '{print NF}' file.txt
```

# Print lines with field count

```
awk '{print NF "fields:", $0}' file.txt
```

# Print first and last fields

```
awk '{print $1, $NF}' file.txt
```

# Swap first and second fields

```
awk '{print $2, $1}' file.txt
```

# Sum first column

```
awk '{sum += $1} END {print sum}' file.txt
```

# Average of first column

```
awk '{sum += $1} END {print sum/NR}' file.txt
```

# Count lines

```
awk 'END {print NR}' file.txt
```

# Count non-empty lines

```
awk 'NF {count++} END {print count}' file.txt
```

# Use comma as separator

```
awk -F',' '{print $1}' file.csv
```

# Use colon as separator

```
awk -F':' '{print $1}' /etc/passwd
```

# Multiple character separator

```
awk -F'[ ,]' '{print $1}' file.txt # Space or comma
```

# Change output separator

```
awk 'BEGIN {OFS=","} {print $1, $2}' file.txt
```



# Print first 10 lines

awk 'NR <= 10' file.txt

# Print lines 5 through 10

awk 'NR >= 5 && NR <= 10' file.txt

# Print odd numbered lines

awk 'NR % 2 == 1' file.txt

# Print even numbered lines

awk 'NR % 2 == 0' file.txt

# Format numbers with spacing  
awk '{printf "%5d %s\n", NR, \$0}' file.txt

# Add zeros before numbers  
awk '{printf "%03d %s\n", NR, \$0}' file.txt

# Format as table  
awk '{printf "%-20s %s\n", \$1, \$2}' file.txt

# Add dashes between fields  
awk '{print \$1 "-" \$2}' file.txt

# Print lines where first field > 100

- `awk '$1 > 100' file.txt`

# # Sum values in third column

- `awk '{sum += $3} END {print sum}' file.txt`

# # Use comma as separator

- `awk -F',' '{print $1}' file.csv`

# Add line numbers

```
awk '{print NR ") " $0}' file.txt
```

# Filter with multiple conditions

*# Print lines where \$1 > 100 and \$2 contains "test"*

```
awk '$1 > 100 && $2 ~ /test/' file.txt
```

# Average of second column

```
awk '{sum += $2} END {print "Average = " sum/NR}' file.txt
```



# Formatted Output

# Format as CSV

- `awk 'BEGIN {OFS=","} {print $1, $2, $3}' file.txt`

# Splitting a Line Into Fields

- `awk '{print $1,$4}' employee.txt`

ajay 45000

sunil 25000

varun 50000

amit 47000

tarun 15000

deepak 23000

sunil 13000

satvik 80000

# Built-In Variables In Awk

- Awk's built-in variables include the field variables—\$1, \$2, \$3, and so on (\$0 is the entire line) — that break a line of text into individual words or pieces called fields.

```
$ awk '{print NR,$0}' employee.txt
```

1 ajay manager account 45000

2 sunil clerk account 25000

3 varun manager sales 50000

4 amit manager account 47000

5 tarun peon sales 15000

6 deepak clerk sales 23000

7 sunil peon sales 13000

8 satvik director purchase 80000

# Use of NF built-in variables (Display Last Field)

```
$ awk '{print $1,$NF}' employee.txt
```

```
ajay 45000  
sunil 25000  
varun 50000  
amit 47000  
tarun 15000  
deepak 23000  
sunil 13000  
satvik 80000
```

```
awk 'NR==3, NR==6 {print NR,$0}' employee.txt
```

3 varun manager sales 50000

4 amit manager account 47000

5 tarun peon sales 15000

6 deepak clerk sales 23000

# SED command

- SED command in UNIX stands for stream editor and it can perform lots of functions on file like searching, find and replace, insertion or deletion.
- SED is a powerful text stream editor.
- Can do insertion, deletion, search and replace(substitution).
- SED command in unix supports regular expression which allows it perform complex pattern matching.

**sed OPTIONS... [SCRIPT] [INPUTFILE...]**

```
$cat > geekfile.txt
```

unix is great os. unix is opensource. unix is free os.

learn operating system.

unix linux which one you choose.

unix is easy to learn.unix is a multiuser os.Learn unix .

unix is a powerful.



# Replacing or substituting string

```
$sed 's/unix/linux/' geekfile.txt
```

linux is great os. unix is opensource.unix is free os.

learn operating system.

linux linux which one you choose.

linux is easy to learn.unix is a multiuser os.Learn unix .unix is a powerful.

“s” specifies the substitution operation.

“/” are delimiters.

# Replacing all the occurrence of the pattern in a line

```
$sed 's/unix/linux/g' geekfile.txt
```

linux is great os. linux is opensource. linux is free os.

learn operating system.

linux linux which one you choose.

linux is easy to learn.linux is a multiuser os.Learn linux .linux is a powerful.

# Replacing from nth occurrence to all occurrences in a line

```
$sed 's/unix/linux/3g' geekfile.txt
```

unix is great os. unix is opensource. linux is free os.

learn operating system.

unix linux which one you choose.

unix is easy to learn.unix is a multiuser os.Learn linux .linux is a powerful.

# Parenthesize first character of each word

```
$ echo "Welcome To The Geek Stuff" | sed 's/\(\b[A-Z]\)/\(\1\)/g'
```

(W)elcome (T)o (T)he (G)eek (S)tuff

# Replacing string on a specific line number

```
$sed '3 s/unix/linux/' geekfile.txt
```

unix is great os. unix is opensource. unix is free os.  
learn operating system.

linux linux which one you choose.

unix is easy to learn.unix is a multiuser os.Learn unix .unix is a powerful.

# Duplicating the replaced line with /p flag

- /p print flag prints the replaced line twice on the terminal. If a line does not have the search pattern and is not replaced, then the /p prints that line only once.

linux is great os. unix is opensource. unix is free os.

linux is great os. unix is opensource. unix is free os.

learn operating system.

linux linux which one you choose.

linux linux which one you choose.

linux is easy to learn.unix is a multiuser os.Learn unix .unix is a powerful.

linux is easy to learn.unix is a multiuser os.Learn unix .unix is a powerful.

# Printing only the replaced lines :

```
$sed -n 's/unix/linux/p' geekfile.txt
```

linux is great os. unix is opensource. unix is free os.

linux linux which one you choose.

linux is easy to learn.unix is a multiuser os.Learn unix .unix is a powerful.

# Replacing string on a range of lines :

- `$sed '1,3 s/unix/linux/' geekfile.txt`

**linux is great os. unix is opensource. unix is free os.**

**learn operating system.**

**linux linux which one you choose.**

unix is easy to learn.unix is a multiuser os.Learn unix .unix is a powerful.



```
$sed '2,$ s/unix/linux/' geekfile.txt
```

unix is great os. unix is opensource. unix is free os.

**learn operating system.**

**linux linux which one you choose.**

**linux is easy to learn.unix is a multiuser os.Learn unix .unix is a powerful**

# Deleting lines from a particular file

To Delete a particular line say n in this example

Syntax:

```
$ sed 'nd' filename.txt
```

Example:

```
$ sed '5d' filename.txt
```

- To Delete a last line

`$ sed '$d' filename.txt`

- To Delete from nth to last line

`$ sed '3,6d' filename.txt`

- To Delete pattern matching line

`$ sed '/abc/d' filename.txt`

cat a.txt

life isn't meant to be easy, life is meant to be lived.

Try to learn & understand something new everyday in life.

Respect everyone & most important love everyone.

Don't hesitate to ask for love & don't hesitate to show love too.

Life is too short to be shy.

In life, experience will help you differentiating right from wrong.

- Insert one blank line after each line : `sed G a.txt`
- To insert two blank lines : `sed 'G;G' a.txt`
- Insert 5 spaces to the left of every lines : `sed 's/^/ /' a.txt`